

# Machine Learning Engineer Capstone Project

## Capstone Proposal

Yemi Awosanya

April 23, 2021

### Table of Contents

## 1. Project Overview

### 1.1 Introduction

### 1.2 Statement

### 1.3 Datasets

## 2. Project Analysis

## 3. Methodology

### 3.1 Data Visualization

### 3.2 Algorithms and Techniques

## 4. Results

## 5. Conclusion

## 6. References

### Project Overview

#### Introduction

This data set contains simulated data that mimics customer behavior on the Starbucks rewards mobile app. Once every few days, Starbucks sends out an offer to users of the mobile app. An offer can be merely an advertisement for a drink or an actual offer such as a discount or BOGO (buy one get one free). Some users might not receive any offer during certain weeks.

Not all users receive the same offer, and that is the challenge to solve with this data set.

Your task is to combine transaction, demographic and offer data to determine which demographic groups respond best to which offer type. This data set is a simplified version of the real Starbucks app because the underlying simulator only has one product whereas Starbucks actually sells dozens of products.

Every offer has a validity period before the offer expires. As an example, a BOGO offer might be valid for only 5 days. You'll see in the data set that informational offers have a validity period even though these ads

are merely providing information about a product; for example, if an informational offer has 7 days of validity, you can assume the customer is feeling the influence of the offer for 7 days after receiving the advertisement.

You'll be given transactional data showing user purchases made on the app including the timestamp of purchase and the amount of money spent on a purchase. This transactional data also has a record for each offer that a user receives as well as a record for when a user actually views the offer. There are also records for when a user completes an offer.

Keep in mind as well that someone using the app might make a purchase through the app without having received an offer or seen an offer.

## Statement

Starbucks is passionate about improving and uplifting the lives of its customers, one beverage at a time. In doing so, they have created a mobile application that rewards and sends offers to their valued customers. The application offers promotions such as discounts or buy-one-get-one-free (BOGO), and occasionally, informational offers.

We used the study of consumer behavior to see how marketing campaigns can be adapted and improved to more effectively influence the consumer.

Because of Starbucks' dedication to their customers, we analyzed customer data to explore ways of improving their experience and to serve them better.

## Datasets

We have three datasets as follows:

### profile.json

#### Rewards program users (17000 users x 5 fields)

- gender: (categorical) M, F, O, or null
- age: (numeric) missing value encoded as 118
- id: (string/hash)
- became\_member\_on: (date) format YYYYMMDD
- income: (numeric)

	gender	age	id	became_member_on	income
0	None	118	68be06ca386d4c31939f3a4f0e3dd783	20170212	NaN
1	F	55	0610b486422d4921ae7d2bf64640c50b	20170715	112000.0
2	None	118	38fe809add3b4fcf9315a9694bb96ff5	20180712	NaN
3	F	75	78afa995795e4d85b5d9ceeca43f5fef	20170509	100000.0
4	None	118	a03223e636434f42ac4c3df47e8bac43	20170804	NaN

**portfolio.json**

Offers sent during 30-day test period (10 offers x 6 fields)

- reward: (numeric) money awarded for the amount spent
- channels: (list) web, email, mobile, social
- difficulty: (numeric) money required to be spent to receive reward
- duration: (numeric) time for offer to be open, in days
- offer\_type: (string) bogo, discount, informational
- id: (string/hash)

	reward	channels	difficulty	duration	offer_type	id	email	v
0	10	[email, mobile, social]	10	7	bogo	ae264e3637204a6fb9bb56bc8210ddfd		1.0
1	10	[web, email, mobile, social]	10	5	bogo	4d5c57ea9a6940dd891ad53e9dbe8da0		1.0
2	0	[web, email, mobile]	0	4	informational	3f207df678b143eea3cee63160fa8bed		1.0
3	5	[web, email, mobile]	5	7	bogo	9b98b8c7a33c4b65b9aebfe6a799e6d9		1.0
4	5	[web, email]	20	10	discount	0b1e1539f2cc45b7b9fa7c272da2e1d7		1.0
5	3	[web, email, mobile, social]	7	7	discount	2298d6c36e964ae4a3e7e9706d1fb8c2		1.0

**transcript.json**

Event log (306648 events x 4 fields)

- person: (string/hash)
- event: (string) offer received, offer viewed, transaction, offer completed
- value: (dictionary) different values depending on event type
- offer id: (string/hash) not associated with any "transaction"
- amount: (numeric) money spent in "transaction"
- reward: (numeric) money gained from "offer completed"
- time: (numeric) hours after start of test

	person	event	value	time
0	78afa995795e4d85b5d9ceeca43f5fef	offer received	{'offer id': '9b98b8c7a33c4b65b9aebfe6a799e6d9'}	0
1	a03223e636434f42ac4c3df47e8bac43	offer received	{'offer id': '0b1e1539f2cc45b7b9fa7c272da2e1d7'}	0
2	e2127556f4f64592b11af22de27a7932	offer received	{'offer id': '2906b810c7d4411798c6938adc9daaa5'}	0
3	8ec6ce2a7e7949b1bf142def7d0e0586	offer received	{'offer id': 'fafdcd668e3743c1bb461111dcafc2a4'}	0
4	68617ca6246f4fbc85e91a2a49552598	offer received	{'offer id': '4d5c57ea9a6940dd891ad53e9dbe8da0'}	0

## Methodology

We followed the steps below to clean and process our datasets:


1. Check dataset content
2. Check data types
3. Check null/empty cells
4. Drop duplicate values, if applicable
5. Fill cell values, if applicable

## Data Visualization


We started with data of over 300000 records. After cleaning the data, we were left with a total record number of 76277.

As part of the process, we plotted a few graphs to gain more insights to the customer data. We

Here, we have the age distrubution which indicates there a normal distribution with some outlier of age, 118.



In the income distribution graph, we can see that the income distribution is not sufficient enough to tell the user' story.



With the datasets, we used the preprocessor from sklearn, and we created new features. One of such feature was to determine and count successful offers within the allotted offer period.

We adopted the 20/80 rule for splitting the training and test data sets.

```
In [46]: dx = df.loc[:, df.columns != 'offer success']
label = df['offer success']

X_train, X_test, y_train, y_test = train_test_split(dx, label, test_size = 0.2, random_state =

print("Training data has {} rows".format(X_train.shape[0]))
print("Testing data has {} rows".format(X_test.shape[0]))

Training data has 61021 rows
Testing data has 15256 rows
```

Algorithms and Techniques

we used the GradientBoostingClassifier as the benchmark model.

	Benchmark Model	Train Score	Test Score
0	GradientBoostingClassifier	0.979925	0.981319

And, we choose xgboost as our final model to predict whether a customer accepts an offer or not.

	Xgboost Model	Train Score	Test Score
0	Xgboost Model	0.995379	0.999148

Results

We used the following metrics to evaluate our model:

- 1. roc\_auc\_score - Area Under the Receiver Operating Characteristic Curve (ROC AUC) from prediction scores
- 2. accuracy\_score - The number of correct predictions made divided by the total number of predictions made
- 3. f1\_score - Weighted average of the precision and recall, where an F1 score reaches its best value at 1

We used the sklearn xgboost tuning algorithm with the following hyperparameters

```
xgboost = XGBClassifier(  
    learning_rate =0.1,  
    n_estimators=1000,  
    max_depth=6,  
    min_child_weight=4,  
    gamma=0,  
    subsample=0.8,  
    colsample_bytree=0.9,  
    reg_alpha=0.005,  
    objective= 'binary:logistic',  
    nthread=3,  
    scale_pos_weight=2,  
    seed=30)
```


## Testing results for the 3 metrics listed above

```
In [51]: xgboost.fit(X_test, y_test)  
  
xgboost_predict = xgboost.predict(X_test)  
  
xgboost_accuracy_score_test = accuracy_score(xgboost_predict, y_test)  
xgboost_roc_auc_score_test = roc_auc_score(xgboost_predict, y_test)  
xgboost_f1_score_test= f1_score(xgboost_predict, y_test, average='macro')  
  
print("Accuracy Score (Test) : %.4g" % xgboost_accuracy_score_test)  
print("AUC Score (Test) : %.4g" % xgboost_roc_auc_score_test)  
print("F1 Score (Test): %.4g" % xgboost_f1_score_test)  
  
Accuracy Score (Test) : 0.9991  
AUC Score (Test) : 0.9782  
F1 Score (Test): 0.9886
```

## Training results for the 3 metrics listed above

```
In [52]: xgboost.fit(X_train, y_train)  
  
xgboost_predict = xgboost.predict(X_train)  
  
xgboost_accuracy_score_train = accuracy_score(xgboost_predict, y_train)  
xgboost_roc_auc_score_train = roc_auc_score(xgboost_predict, y_train)  
xgboost_f1_score_train= f1_score(xgboost_predict, y_train, average='macro')  
  
print("Accuracy Score (Train) : %.4g" % xgboost_accuracy_score_train)  
print("AUC Score (Train): %f" % xgboost_roc_auc_score_train)  
print("F1 Score (Train): %f" % xgboost_f1_score_train)  
  
Accuracy Score (Train) : 0.9954  
AUC Score (Train): 0.912239  
F1 Score (Train): 0.946128
```

## Conclusion



True positive of 59542, False positive of 28, False negative of 254, True negative of 1197, it fair to assume the model predicts whether a customer will accept an offer or not, accurately.

## References

```
import matplotlib.pyplot as plt
import seaborn as sns

import pandas as pd
import numpy as np
%matplotlib inline

from sklearn.preprocessing import scale
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split

from sklearn import preprocessing

from datetime import datetime
from datetime import date
from dateutil.relativedelta import relativedelta

from xgboost import XGBClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.metrics import accuracy_score, roc_auc_score, f1_score, confusion_matrix, auc
```

- [Feature Selection](#)



- [Evaluate Gradient Boosting Models](#)
- [Getting Started with XGBoost](#)
- [Categorical Data Encoding](#)